

THE LINUX



COMMANDS HANDBOOK

Flavio Copes

passwd
ping
traceroute
clear
history
export
crontab
uname
env
printenv
Conclusion

Table of Contents

Preface
Introduction to Linux and shells
man
ls
cd
pwd
mkdir
rmdir
mv
cp
open
touch
find
ln
gzip
gunzip
tar
alias
cat
less
tail
wc
grep
sort

uniq
diff
echo
chown
chmod
umask
du
df
basename
dirname
ps
top
kill
killall
jobs
bg
fg
type
which
nohup
xargs
vim
emacs
nano
whoami
who
su
sudo

1

2

Preface

The Linux Commands Handbook follows the 80/20 rule: learn in 20% of the time the 80% of a topic.

I find this approach gives a well-rounded overview.

This book does not try to cover everything under the sun related to Linux and its commands. It focuses on the small core commands that you will use the 80% or 90% of the time, trying to simplify the usage of the more complex ones.

All those commands work on Linux, macOS, WSL, and anywhere you have a UNIX environment.

I hope the contents of this book will help you achieve what you want: **get comfortable with Linux**.

This book is written by Flavio. I **publish programming tutorials** every day on my website flaviocopes.com.

You can reach me on Twitter [@flaviocopes](https://twitter.com/flaviocopes).

Enjoy!

Introduction to Linux and shells

Linux is an operating system, like macOS or Windows.

It is also the most popular Open Source and free, as in freedom, operating system.

It powers the vast majority of the servers that compose the Internet. It's the base upon which everything is built upon. But not just that. Android is based on (a modified version of) Linux.

The Linux "core" (called *kernel*) was born in 1991 in Finland, and it went a really long way from its humble beginnings. It went on to be the kernel of the GNU Operating System, creating the duo GNU/Linux.

There's one thing about Linux that corporations like Microsoft and Apple, or Google, will never be able to offer: the freedom to do whatever you want with your computer.

They're actually going in the opposite direction, building walled gardens, especially on the mobile side.

Linux is the ultimate freedom.

It is developed by volunteers, some paid by companies that rely on it, some independently, but there's no single commercial company that can dictate what goes into Linux, or the project priorities.

Linux can also be used as your day to day computer. I use macOS because I really enjoy the applications, the design and I also used to be an iOS and Mac apps developer, but before using it I used Linux as my main computer Operating System.

No one can dictate which apps you can run, or "call home" with apps that track you, your position, and more.

Linux is also special because there's not just "one Linux", like it happens on Windows or macOS. Instead, we have **distributions**.

A "distro" is made by a company or organization and packages the Linux core with additional programs and tooling.

For example you have Debian, Red Hat, and Ubuntu, probably the most popular.

Many, many more exist. You can create your own distribution, too. But most likely you'll use a popular one, one that has lots of users and a community of people around it, so you can do what you need to do without losing too much time reinventing the wheel and figuring out answers to common problems.

Some desktop computers and laptops ship with Linux preinstalled. Or you can install it on your Windows-based computer, or on a Mac.

But you don't need to disrupt your existing computer just to get an idea of how Linux works.

I don't have a Linux computer.

If you use a Mac you need to know that under the hood macOS is a UNIX Operating System, and it shares a lot of the same ideas and software that a GNU/Linux system uses, because GNU/Linux is a free alternative to UNIX.

UNIX is an umbrella term that groups many operating systems used in big corporations and institutions, starting from the 70's

The macOS terminal gives you access to the same exact commands I'll describe in the rest of this handbook.

Microsoft has an official [Windows Subsystem for Linux](#) which you can (and should!) install on Windows. This will give you the ability to run Linux in a very easy way on your PC.

But the vast majority of the time you will run a Linux computer in the cloud via a VPS (Virtual Private Server) like DigitalOcean.

A shell is a command interpreter that exposes to the user an interface to work with the underlying operating system.

It allows you to execute operations using text and commands, and it provides users advanced features like being able to create scripts.

This is important: shells let you perform things in a more optimized way than a GUI (Graphical User Interface) could ever possibly let you do. Command line tools can offer many different configuration options without being too complex to use.

There are many different kind of shells. This post focuses on Unix shells, the ones that you will find commonly on Linux and macOS computers.

Many different kind of shells were created for those systems over time, and a few of them dominate the space: Bash, Csh, Zsh, Fish and many more!

All shells originate from the Bourne Shell, called `sh`. "Bourne" because its creator was Steve Bourne.

Bash means *Bourne-again shell*. `sh` was proprietary and not open source, and Bash was created in 1989 to create a free alternative for the GNU project and the Free Software Foundation. Since projects had to pay to use the Bourne shell, Bash became very popular.

If you use a Mac, try opening your Mac terminal. That by default is running ZSH. (or, pre-Catalina, Bash)

You can set up your system to run any kind of shell, for example I use the Fish shell.

Each single shell has its own unique features and advanced usage, but they all share a common functionality: they can let you execute programs, and they can be programmed.

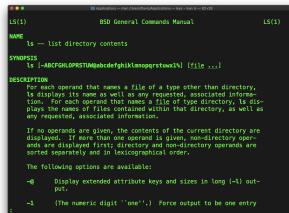
In the rest of this handbook we'll see in detail the most common commands you will use.

6

man

The first command I want to introduce is a command that will help you understand all the other commands.

Every time I don't know how to use a command, I type `man <command>` to get the manual:



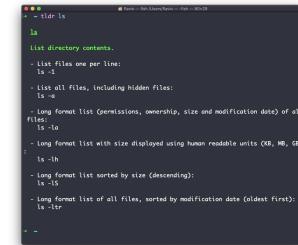
The screenshot shows a terminal window with the title "BSD General Commands Manual". It displays the man page for the "ls" command. The page starts with a brief description of the command and then lists several options: -l, -1, -o, -L, -h, -S, -t, -r, and -v. Each option is followed by a detailed explanation of its purpose and behavior.

This is a man (from *manual*) page. Man pages are an essential tool to learn, as a developer. They contain so much information that sometimes it's almost too much.

The above screenshot is just 1 of 14 screens of explanation for the `ls` command.

Most of the times when I'm in need to learn a command quickly I use this site called [tldr pages](#): <https://tldr.sh/>. It's a command you can install, then you run it like this: `tldr <command>`, which gives you a very quick overview of a command, with some handy examples of common usage scenarios:

7



The screenshot shows a terminal window with the title "tldr ls". It displays the output of the tldr command for the ls command. The output is a brief summary of the command, listing its name, synopsis, and a detailed description of its various options: -l, -1, -o, -L, -h, -S, -t, -r, and -v. Each option is followed by a short description of its function.

This is not a substitute for `man`, but a handy tool to avoid losing yourself in the huge amount of information present in a man page. Then you can use the `man` page to explore all the different options and parameters you can use on a command.

8

ls

Inside a folder you can list all the files that the folder contains using the `ls` command:

```
ls
```

If you add a folder name or path, it will print that folder contents:

```
ls /bin
```



The screenshot shows a terminal window with the title "Terminal". It displays the output of the command `ls /bin`. The output lists several executable files located in the /bin directory, including bash, csh, ed, launchctl, mv, rm, rmtrash, tcsh, cat, dd, hostname, ln, ls, sleep, unlink, chmod, df, kill, lsblk, pad, sync, wallpath, cp, echo, rm, skdirm, rsync, and zsh.

`ls` accepts a lot of options. One of my favorite options combinations is `-al`. Try it:

```
ls -al /bin
```

9

10

11

```
ls -al /bin
total 5120
drwxr-xr-x 37 root wheel 114k Feb  4 18:09 .
drwxr-xr-x  1 root wheel 22784 Jan 16 02:21 [
drwxr-xr-x  1 root wheel 1048k Jan 16 02:21 bin
drwxr-xr-x  1 root wheel 23648 Jan 16 02:21 cat
drwxr-xr-x  1 root wheel 1536k Jan 16 02:21 cron
drwxr-xr-x  1 root wheel 29824 Jan 16 02:21 cp
drwxr-xr-x  1 root wheel 379952 Jan 16 02:21 csh
drwxr-xr-x  1 root wheel 1048k Jan 16 02:21 cpio
drwxr-xr-x  1 root wheel 32480 Jan 16 02:21 dd
drwxr-xr-x  1 root wheel 33520 Jan 16 02:21 diff
drwxr-xr-x  1 root wheel 18128 Jan 16 02:21 echo
drwxr-xr-x  1 root wheel 2480k Jan 16 02:21 expr
drwxr-xr-x  1 root wheel 23192 Jan 16 02:21 grep
drwxr-xr-x  1 root wheel 18288 Jan 16 02:21 hostname
drwxr-xr-x  1 root wheel 1048k Jan 16 02:21 kill
drwxr-xr-x  1 root wheel 128264 Jan 16 02:21 ksh
drwxr-xr-x  1 root wheel 123196 Jan 16 02:21 launchctl
```

compared to the plain `ls`, this returns much more information.

You have, from left to right:

- the file permissions (and if your system supports ACLs, you get an ACL flag as well)
- the number of links to that file
- the owner of the file
- the group of the file
- the file size in bytes
- the file modified datetime
- the file name

This set of data is generated by the `l` option. The `a` option instead also shows the hidden files.

Hidden files are files that start with a dot (`.`).

cd

Once you have a folder, you can move into it using the `cd` command. `cd` means change directory. You invoke it specifying a folder to move into. You can specify a folder name, or an entire path.

Example:

```
mkdir fruits
cd fruits
```

Now you are into the `fruits` folder.

You can use the `..` special path to indicate the parent folder:

```
cd .. #back to the home folder
```

The `#` character indicates the start of the comment, which lasts for the entire line after it's found.

You can use it to form a path:

```
mkdir fruits
mkdir cars
cd fruits
cd ../cars
```

There is another special path indicator which is `..` and indicates the **current** folder.

You can also use absolute paths, which start from the root folder `/`:

`cd /etc`

This command works on Linux, macOS, WSL, and anywhere you have a UNIX environment

pwd

Whenever you feel lost in the filesystem, call the `pwd` command to know where you are:

```
pwd
```

It will print the current folder path.

mkdir

You create folders using the `mkdir` command:

```
mkdir fruits
```

You can create multiple folders with one command:

```
mkdir dogs cars
```

You can also create multiple nested folders by adding the `-p` option:

```
mkdir -p fruits/apples
```

Options in UNIX commands commonly take this form. You add them right after the command name, and they change how the command behaves. You can often combine multiple options, too.

You can find which options a command supports by typing `man <commandname>`. Try now with `man mkdir` for example (press the `q` key to esc the man page). Man pages are the amazing built-in help for UNIX.

rmdir

Just as you can create a folder using `mkdir`, you can delete a folder using `rmdir`:

```
rmdir fruits
rmdir fruits
```

You can also delete multiple folders at once:

```
mkdir fruits cars
rmdir fruits cars
```

The folder you delete must be empty.

To delete folders with files in them, we'll use the more generic `rm` command which deletes files and folders, using the `-rf` options:

```
rm -rf fruits cars
```

Be careful as this command does not ask for confirmation and it will immediately remove anything you ask it to remove.

There is no `bin` when removing files from the command line, and recovering lost files can be hard.

mv

Once you have a file, you can move it around using the `mv` command. You specify the file current path, and its new path:

```
touch test  
mv pear new_pear
```

The `pear` file is now moved to `new_pear`. This is how you **rename** files and folders.

If the last parameter is a folder, the file located at the first parameter path is going to be moved into that folder. In this case, you can specify a list of files and they will all be moved in the folder path identified by the last parameter:

```
touch pear  
touch apple  
mkdir fruits  
mv pear apple fruits #pear and apple moved to the fi
```

cp

You can copy a file using the `cp` command:

```
touch test  
cp apple another_apple
```

To copy folders you need to add the `-r` option to recursively copy the whole folder contents:

```
mkdir fruits  
cp -r fruits cars
```

open

The `open` command lets you open a file using this syntax:

```
open <filename>
```

You can also open a directory, which on macOS opens the Finder app with the current directory open:

```
open <directory name>
```

I use it all the time to open the current directory:

```
open .
```

The special `.` symbol points to the current directory, as `..` points to the parent directory

The same command can also be used to run an application:

```
open <application name>
```

18

19

20

touch

You can create an empty file using the `touch` command:

```
touch apple
```

If the file already exists, it opens the file in write mode, and the timestamp of the file is updated.

find

The `find` command can be used to find files or folders matching a particular search pattern. It searches recursively.

Let's learn it by example.

Find all the files under the current tree that have the `.js` extension and print the relative path of each file matching:

```
find . -name '*.js'
```

It's important to use quotes around special characters like `*` to avoid the shell interpreting them.

Find directories under the current tree matching the name "src":

```
find . -type d -name src
```

Use `-type f` to search only files, or `-type l` to only search symbolic links.

`-name` is case sensitive. Use `-iname` to perform a case-insensitive search.

You can search under multiple root trees:

```
find folder1 folder2 -name filename.txt
```

Find directories under the current tree matching the name "node_modules" or "public":

```
find . -type d -name node_modules -or -name public
```

You can also exclude a path, using `-not -path`:

```
find . -type d -name '*.*md' -not -path 'node_modules'
```

You can search files that have more than 100 characters (bytes) in them:

```
find . -type f -size +100c
```

Search files bigger than 100KB but smaller than 1MB:

```
find . -type f -size +100k -size -1M
```

Search files edited more than 3 days ago

```
find . -type f -mtime +3
```

Search files edited in the last 24 hours

```
find . -type f -mtime -1
```

You can delete all the files matching a search by adding the `-delete` option. This deletes all the files edited in the last 24 hours:

```
find . -type f -mtime -1 -delete
```

21

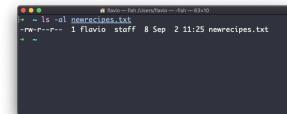
22

23

You can execute a command on each result of the search. In this example we run `cat` to print the file content:

```
find . -type f -exec cat {} \;
```

notice the terminating `\; . .` is filled with the file name at execution time.



```
ls -al newrecipes.txt
-rw-r--r-- 1 flavio staff 8 Sep 2 11:25 newrecipes.txt
```

In

The `ln` command is part of the Linux file system commands.

It's used to create links. What is a link? It's like a pointer to another file. A file that points to another file. You might be familiar with Windows shortcuts. They're similar.

We have 2 types of links: **hard links** and **soft links**.

Hard links

Hard links are rarely used. They have a few limitations: you can't link to directories, and you can't link to external filesystems (disks).

A hard link is created using

```
ln <original> <link>
```

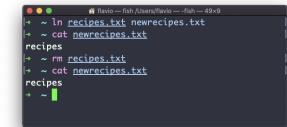
For example, say you have a file called `recipes.txt`. You can create a hard link to it using:

```
ln recipes.txt newrecipes.txt
```

The new hard link you created is indistinguishable from a regular file:

Now any time you edit any of those files, the content will be updated for both.

If you delete the original file, the link will still contain the original file content, as that's not removed until there is one hard link pointing to it.



```
ln recipes.txt newrecipes.txt
rm recipes.txt
```

Soft links

Soft links are different. They are more powerful as you can link to other filesystems and to directories, but when the original is removed, the link will be broken.

You create soft links using the `-s` option of `ln`:

24

25

26

```
ln -s <original> <link>
```

For example, say you have a file called `recipes.txt`. You can create a soft link to it using:

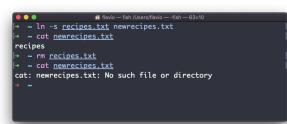
```
ln -s recipes.txt newrecipes.txt
```

In this case you can see there's a special `l` flag when you list the file using `ls -al`, and the file name has a `@` at the end, and it's colored differently if you have colors enabled:



```
ln -s newrecipes.txt newrecipes.txt
rm newrecipes.txt
```

Now if you delete the original file, the links will be broken, and the shell will tell you "No such file or directory" if you try to access it:



```
ln -s recipes.txt newrecipes.txt
rm recipes.txt
ls -al newrecipes.txt
newrecipes.txt: No such file or directory
```

gzip

You can compress a file using the gzip compression protocol named LZ77 using the `gzip` command.

Here's the simplest usage:

```
gzip filename
```

This will compress the file, and append a `.gz` extension to it. The original file is deleted. To prevent this, you can use the `-c` option and use output redirection to write the output to the `filename.gz` file:

```
gzip -c filename > filename.gz
```

The `-c` option specifies that output will go to the standard output stream, leaving the original file intact

Or you can use the `-k` option:

```
gzip -k filename
```

There are various levels of compression. The more the compression, the longer it will take to compress (and decompress). Levels range from 1 (fastest, worst compression) to 9 (slowest, better compression), and the default is 6.

You can choose a specific level with the `--<NUMBER>` option:

```
gzip -1 filename
```

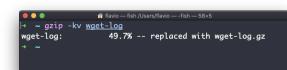
You can compress multiple files by listing them:

```
gzip filename1 filename2
```

You can compress all the files in a directory, recursively, using the `-r` option:

```
gzip -r a_folder
```

The `-v` option prints the compression percentage information. Here's an example of it being used along with the `-k` (keep) option:



```
gzip -kv a_folder
a_folder.gz: 49.7% -- replaced with a_folder.log
```

`gzip` can also be used to decompress a file, using the `-d` option:

```
gzip -d filename.gz
```

27

28

29

gunzip

The `gunzip` command is basically equivalent to the `gzip` command, except the `-d` option is always enabled by default.

The command can be invoked in this way:

```
gunzip filename.gz
```

This will gunzip and will remove the `.gz` extension, putting the result in the `filename` file. If that file exists, it will overwrite that.

You can extract to a different filename using output redirection using the `-c` option:

```
gunzip -c filename.gz > anotherfilename
```

30

tar

The `tar` command is used to create an archive, grouping multiple files in a single file.

Its name comes from the past and means *tape archive*. Back when archives were stored on tapes.

This command creates an archive named `archive.tar` with the content of `file1` and `file2`:

```
tar -cf archive.tar file1 file2
```

The `c` option stands for *create*. The `f` option is used to write to file the archive.

To extract files from an archive in the current folder, use:

```
tar -xf archive.tar
```

the `x` option stands for *extract*

and to extract them to a specific directory, use:

```
tar -xvf archive.tar -C directory
```

You can also just list the files contained in an archive:



`tar` is often used to create a **compressed archive**, gzipping the archive.

This is done using the `z` option:

```
tar -czf archive.tar.gz file1 file2
```

This is just like creating a tar archive, and then running `gzip` on it.

To unarchive a gzipped archive, you can use `gunzip`, or `gzip -d`, and then unarchive it, but `tar -xf` will recognize it's a gzipped archive, and do it for you:

```
tar -xf archive.tar.gz
```

32

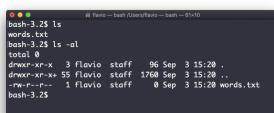
alias

It's common to always run a program with a set of options you like using.

For example, take the `ls` command. By default it prints very little information:



while using the `-al` option it will print something more useful, including the file modification date, the size, the owner, and the permissions, also listing hidden files (files starting with a `.`):

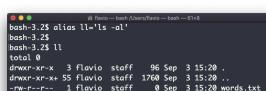


You can create a new command, for example I like to call it `ll`, that is an alias to `ls -al`.

You do it in this way:

```
alias ll='ls -al'
```

Once you do, you can call `ll` just like it was a regular UNIX command:



Now calling `alias` without any option will list the aliases defined:



The alias will work until the terminal session is closed.

To make it permanent, you need to add it to the shell configuration, which could be `~/.bashrc` or `~/.profile` or `~/.bash_profile` if you use the Bash shell, depending on the use case.

Be careful with quotes if you have variables in the command: using double quotes the variable is resolved at definition time, using single quotes it's resolved at invocation time. Those 2 are different:

```
alias lsthis="ls $PWD"  
alias lscurrent='ls $PWD'
```

`SPWD` refers to the current folder the shell is into. If you now navigate away to a new folder, `lscurrent` lists the files in the new folder, `lsthis` still lists the files in the folder you were when you defined the alias.

33

34

35

cat

Similar to `tail` in some way, we have `cat`. Except `cat` can also add content to a file, and this makes it super powerful.

In its simplest usage, `cat` prints a file's content to the standard output:

```
cat file
```

You can print the content of multiple files:

```
cat file1 file2
```

and using the output redirection operator `>` you can concatenate the content of multiple files into a new file:

```
cat file1 file2 > file3
```

Using `>>` you can append the content of multiple files into a new file, creating it if it does not exist:

```
cat file1 file2 >> file3
```

When watching source code files it's great to see the line numbers, and you can have `cat` print them using the `-n` option:

```
cat -n file1
```

36

You can only add a number to non-blank lines using `-b`, or you can also remove all the multiple empty lines using `-s`.

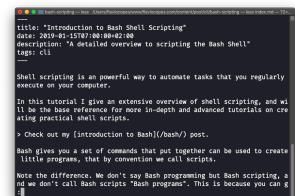
`cat` is often used in combination with the pipe operator `|` to feed a file content as input to another command: `cat file1 | anothercommand`.

40

less

The `less` command is one I use a lot. It shows you the content stored inside a file, in a nice and interactive UI.

Usage: `less <filename>`.



```
title: "Introduction to Bash Shell Scripting"
date: 2019-03-15T07:00:00Z
tags: cli
-----
Shell scripting is an powerful way to automate tasks that you regularly execute on your computer.
In this tutorial I give an extensive overview of shell scripting, and will be covering more in-depth and advanced tutorials on creating practical shell scripts.
> Check out my [Introduction to Bash](/bash/) post.
Bash gives you a set of commands that put together can be used to create little programs, that by convention we call scripts.
Note the difference. We don't say Bash programming but Bash scripting, and we don't call Bash scripts "Bash programs". This is because you can g
```

Once you are inside a `less` session, you can quit by pressing `q`.

You can navigate the file contents using the `up` and `down` keys, or using the `space bar` and `b` to navigate page by page. You can also jump to the end of the file pressing `G` and jump back to the start pressing `g`.

You can search contents inside the file by pressing `/` and typing a word to search. This searches *forward*. You can search backwards using the `?` symbol and typing a word.

38

tail

The best use case of `tail` in my opinion is when called with the `-f` option. It opens the file at the end, and watches for file changes. Any time there is new content in the file, it is printed in the window. This is great for watching log files, for example:

```
tail -f /var/log/system.log
```

To exit, press `ctrl-C`.

You can print the last 10 lines in a file:

```
tail -n 10 <filename>
```

You can print the whole file content starting from a specific line using `+` before the line number:

```
tail -n +10 <filename>
```

`tail` can do much more and as always my advice is to check `man tail`.

41

wc

The `wc` command gives us useful information about a file or input it receives via pipes.

```
echo test >> test.txt
wc test.txt
1 1 5 test.txt
```

Example via pipes, we can count the output of running the `ls -al` command:

```
ls -al | wc
6 47 284
```

The first column returned is the number of lines. The second is the number of words. The third is the number of bytes.

We can tell it to just count the lines:

```
wc -l test.txt
```

or just the words:

```
wc -w test.txt
```

or just the bytes:

```
wc -c test.txt
```

Bytes in ASCII charsets equate to characters, but with non-ASCII charsets, the number of characters might differ because some characters might take multiple bytes, for example this happens in Unicode.

In this case the `-m` flag will help getting the correct value:

```
wc -m test.txt
```

grep

The `grep` command is a very useful tool, that when you master will help you tremendously in your day to day.

If you're wondering, `grep` stands for *global regular expression print*

You can use `grep` to search in files, or combine it with pipes to filter the output of another command.

For example here's how we can find the occurrences of the `document.getElementById` line in the `index.md` file:

```
grep document.getElementById index.md
```

```
bussh-3:25 grep document.getElementById index.md
bussh-3:25
```

Using the `-n` option it will show the line numbers:

```
grep -n document.getElementById index.md
```

```
bussh-3:25 grep -n document.getElementById index.md
bussh-3:25
```

One very useful thing is to tell `grep` to print 2 lines before, and 2 lines after the matched line, to give us more context. That's done using the `-C` option, which accepts a number of lines:

```
grep -nC 2 document.getElementById index.md
```

```
bussh-3:25 grep -nC 2 document.getElementById index.md
bussh-3:25
```

Search is case sensitive by default. Use the `-i` flag to make it insensitive.

As mentioned, you can use `grep` to filter the output of another command. We can replicate the same functionality as above using:

```
less index.md | grep -n document.getElementById
```

```
bussh-3:25 less index.md | grep -n document.getElementById
bussh-3:25
```

The search string can be a regular expression, and this makes `grep` very powerful.

42

43

44

Another thing you might find very useful is to invert the result, excluding the lines that match a particular string, using the `-v` option:

```
bussh-3:25 ls -al
total 128
drwxr-xr-x  4 Florio  staff  128 Apr  3 08:15 .
drwxr-xr-x 138 Florio  staff  4416 Apr 22 14:53 ..
-rw-r--r--  1 Florio  staff  49556 Apr  3 09:37 banner.png
-rw-r--r--  1 Florio  staff  5659 Feb 11 2020 index.md
bussh-3:25 ls -al | grep -v index.md
total 128
drwxr-xr-x  4 Florio  staff  128 Apr  3 08:15 .
drwxr-xr-x 138 Florio  staff  4416 Apr 22 14:53 ..
-rw-r--r--  1 Florio  staff  49556 Apr  3 09:37 banner.png
bussh-3:25
```

sort

Suppose you have a text file which contains the names of dogs:

```
bussh-3:25 cat dogs.txt
Roger
Syd
Vanille
Luna
Ivica
Tina
...
[END]
```

This list is unordered.

The `sort` command helps us sorting them by name:

```
bussh-3:25 sort dogs.txt
Ivica
Luna
Roger
Syd
Tina
Vanille
[END]
```

Use the `r` option to reverse the order:

Sorting by default is case sensitive, and alphabetic. Use the `--ignore-case` option to sort case insensitive, and the `-n` option to sort using a numeric order.

If the file contains duplicate lines:

```
GNU nano 2.8.6  File: dogs.txt  Modified:
Roger
Syd
Vanille
Luna
Ivica
Tina
Roger
Syd
[Cancelled]
X Get Help X Write & R Read F1 Y Prev Pw X Cut Text C Cur Pos
X Exit X Justify X Where Iw Y Next Pg X Uncut T1 X To Spell
```

You can use the `-u` option to remove them:

45

46

47

```
Flavio - fish /Users/flavio - fish - 40x8
+ ~ sort -u dogs.txt
Ivica
Luna
Roger
Syd
Tina
Vanille
+ ~
```

`sort` does not just work on files, as many UNIX commands it also works with pipes, so you can use on the output of another command, for example you can order the files returned by `ls` with:

```
ls | sort
```

`sort` is very powerful and has lots more options, which you can explore calling `man sort`.

unid

`uniq` is a command useful to sort lines of text.

You can get those lines from a file, or using pipes from the output of another command:

```
uniq dogs.txt
```

You need to consider this key thing: `uniq` will only detect adjacent duplicate lines.

This implies that you will most likely use it along with `sort`:

```
sort dogs.txt | uni
```

The `sort` command has its own way to remove duplicates with the `-u (unique)` option. But `uniq` has more power.

By default it removes duplicate lines:

48

4

```
λ ~ llavio — fish /Users/llavio — fish - 40x17
~ cat dogs.txt
Roger
Syd
Vanille
Luna
Ivica
Tina
Roger
Syd
` ~ sort dogs.txt | uniq
Ivica
Luna
Roger
Syd
Tina
Vanille
→ ~
```

You can tell it to only display duplicate lines, for example with the `-d` option:

```
sort dogs.txt | uniq -d
```

```
flavio -> fish /Users/flavio --fish -- -40x13
~ - cat dogs.txt
Roger
Syd
Vanille
Luna
Ivica
Tina
Roger
Syd
~ - sort dogs.txt | uniq -d
Roger
Syd
~ -
```

You can use the `-u` option to only display non-duplicate lines:

```
[  ~  cat dogs.txt
Roger
Syd
Vanille
Luna
Ivica
Tina
Roger
Syd
[  ~  sort dogs.txt | uniq -u
Ivica
Luna
Tina
Vanille
[  ~  ]
```

You can count the occurrences of each line with the `c` option:

```
[... ~] $ flavio - fish /Users/flavio -- fish - 40x17
[... ~] $ cat dogs.txt
Roger
Syd
Vanille
Luna
Ivica
Tina
Roger
Syd
[+ ~] $ sort dogs.txt | uniq -c
 1 Ivica
 1 Luna
 2 Roger
 2 Syd
 1 Tina
 1 Vanille
[+ ~] $
```

Use the special combination

```
sort -dacs.txt | uniq -c | sort -pr
```

to then sort these lines by most frequent:

```
Flavio - fish /Users/Flavio -- fish - 44> sort dogs.txt | uniq -c | sort -nr
+ ~ sort dogs.txt | uniq -c | sort -nr
2 Syd
2 Roger
1 Vanille
1 Tina
1 Luna
1 Ivica
```

51

6


```
bash-3.2$ echo $!ls -al
total 8
drwxr-xr-x  2 flavio staff 128 Sep 3 15:43 .
drwxr-xr-x  4 flavio staff 1768 Sep 3 15:20 ..
-rw-r--r--  1 flavio staff  6 Sep 3 15:43 output.txt
-rw-r--r--  1 flavio staff  6 Sep 3 15:43 words.txt
bash-3.2$
```

Note that whitespace is not preserved by default. You need to wrap the command in double quotes to do so:

```
bash-3.2$ echo *ls -al
total 8
drwxr-xr-x  4 flavio staff 128 Sep 3 15:43 .
drwxr-xr-x  55 flavio staff 1768 Sep 3 15:20 ..
-rw-r--r--  1 flavio staff  6 Sep 3 15:43 output.txt
-rw-r--r--  1 flavio staff  6 Sep 3 15:43 words.txt
bash-3.2$
```

You can generate a list of strings, for example ranges:

```
echo {1..5}
```

```
bash-3.2$ echo {1..5}
1 2 3 4
5
bash-3.2$
```

chown

Every file/directory in an Operating System like Linux or macOS (and every UNIX systems in general) has an **owner**.

The owner of a file can do everything with it. It can decide the fate of that file.

The owner (and the **root** user) can change the owner to another user, too, using the **chown** command:

```
chown <owner> <file>
```

Like this:

```
chown flavio test.txt
```

For example if you have a file that's owned by **root**, you can't write to it as another user:

```
bash-3.2$ sudo touch test.txt
bash-3.2$ echo test >> test.txt
bash-3.2$ open test.txt
bash-3.2$
```

You can use **chown** to transfer the ownership to you:

```
sudo touch test.txt
echo test >> test.txt
doh: file error while redirecting file 'test.txt'
open: Permission denied
sudo chown flavio test.txt
echo test >> test.txt
=
```

It's rather common to have the need to change the ownership of a directory, and recursively all the files contained, plus all the subdirectories and the files contained in them, too.

You can do so using the **-R** flag:

```
chown -R <owner> <file>
```

Files/directories don't just have an owner, they also have a **group**. Through this command you can change that simultaneously while you change the owner:

```
chown <owner>:<group> <file>
```

Example:

```
chown flavio:users test.txt
```

You can also just change the group of a file using the **chgrp** command:

```
chgrp <group> <filename>
```

chmod

Every file in the Linux / macOS Operating Systems (and UNIX systems in general) has 3 permissions: Read, write, execute.

Go into a folder, and run the **ls -al** command.

```
bash-3.2$ git(bittrecker) ~ ls -al
total 1576
drwxr-xr-x  2 flavio staff 128 Feb 17 2020 /
drwxr-xr-x  4 flavio staff 448 Apr 29 17:11 .
drwxr-xr-x  4 flavio staff 448 Sep 24 2018 .XStore
drwxr-xr-x  4 flavio staff 448 Sep 24 2018 .XStore.lock
drwxr-xr-x  1 flavio staff 214 Jun 23 2018 gitignore
drwxr-xr-x  1 flavio staff 214 Jun 23 2018 package-lock.json
drwxr-xr-x  1 flavio staff 485968 Jul 1 2018 package-lock.jsonn
drwxr-xr-x  1 flavio staff 1083 Jul 1 2018 package.json
drwxr-xr-x  1 flavio staff 1083 Jul 1 2018 package-lock.jsonn
drwxr-xr-x  5 flavio staff 168 Jun 24 2018 public/
drwxr-xr-x  1 flavio staff 27398 Jun 24 2018 tailwind.js
drwxr-xr-x  1 flavio staff 269363 Jun 24 2018 yarn.lock
bittrecker git(bittrecker) ~
```

The weird strings you see on each file line, like **drwxr-x-**, define the permissions of the file or folder.

Let's dissect it.

The first letter indicates the type of file:

- **-** means it's a normal file
- **d** means it's a directory
- **l** means it's a link

Then you have 3 sets of values:

- The first set represents the permissions of the **owner** of the file
- The second set represents the permissions of the members of the **group** the file is associated to

- The third set represents the permissions of the **everyone else**

Those sets are composed by 3 values. **rwx** means that specific **persona** has read, write and execution access. Anything that is removed is swapped with a **-**, which lets you form various combinations of values and relative permissions: **rwx**, **r--**, **r-x**, and so on.

You can change the permissions given to a file using the **chmod** command.

chmod can be used in 2 ways. The first is using symbolic arguments, the second is using numeric arguments. Let's start with symbols first, which is more intuitive.

You type **chmod** followed by a space, and a letter:

- **a** stands for **all**
- **u** stands for **user**
- **g** stands for **group**
- **o** stands for **others**

Then you type either **+** or **-** to add a permission, or to remove it. Then you enter one or more permissions symbols (**r**, **w**, **x**).

All followed by the file or folder name.

Here are some examples:

```
chmod a+r filename #everyone can now read
chmod a+rwx filename #everyone can now read and write
chmod o-rwx filename #others (not the owner, not in
```

You can apply the same permissions to multiple personas by adding multiple letters before the **+ / -**:

```
chmod og-r filename #other and group can't read any
```

In case you are editing a folder, you can apply the permissions to every file contained in that folder using the **-r** (recursive) flag.

Numeric arguments are faster but I find them hard to remember when you are not using them day to day. You use a digit that represents the permissions of the persona. This number value can be a maximum of 7, and it's calculated in this way:

- 1 if has execution permission
- 2 if has write permission
- 4 if has read permission

This gives us 4 combinations:

- 0 no permissions
- 1 can execute
- 2 can write
- 3 can write, execute
- 4 can read
- 5 can read, execute
- 6 can read, write
- 7 can read, write and execute

We use them in pairs of 3, to set the permissions of all the 3 groups altogether:

```
chmod 777 filename  
chmod 755 filename  
chmod 644 filename
```

umask

When you create a file, you don't have to decide permissions up front. Permissions have defaults.

Those defaults can be controlled and modified using the `umask` command.

Typing `umask` with no arguments will show you the current umask, in this case `0022`:

```
flavio@flavio-laptop:~$ umask  
0022  
flavio@flavio-laptop:~$
```

What does `0022` mean? That's an octal value that represent the permissions.

Another common value is `0002`.

Use `umask -S` to see a human-readable notation:

```
flavio@flavio-laptop:~$ umask -S  
0022  
flavio@flavio-laptop:~$ umask -S  
u=rwx,g=rwx,o=rwx  
flavio@flavio-laptop:~$
```

In this case, the user (`u`), owner of the file, has read, write and execution permissions on files.

Other users belonging to the same group (`g`) have read and execution permission, same as all the other users (`o`).

In the numeric notation, we typically change the last 3 digits.

Here's a list that gives a meaning to the number:

- 0 read, write, execute
- 1 read and write
- 2 read and execute
- 3 read only
- 4 write and execute
- 5 write only
- 6 execute only
- 7 no permissions

Note that this numeric notation differs from the one we use in `chmod`.

We can set a new value for the mask setting the value in numeric format:

```
umask 002
```

or you can change a specific role's permission:

```
umask g+r
```

66

67

68

du

The `du` command will calculate the size of a directory as a whole:

```
du
```

```
flavio@flavio-laptop:~$ du  
32  
flavio@flavio-laptop:~$
```

The `32` number here is a value expressed in bytes.

Running `du *` will calculate the size of each file individually:

```
flavio@flavio-laptop:~$ du *\n8 Card.vue\n8 CardList.vue\n8 Form.vue\n8 HelloWorld.vue\nflavio@flavio-laptop:~$
```

You can set `du` to display values in MegaBytes using `du -m`, and GigaBytes using `du -g`.

The `-h` option will show a human-readable notation for sizes, adapting to the size:

```
flavio@flavio-laptop:~$ du -h vuehandbook\n4.0K vuehandbook/12-vue-watchers\n4.0K vuehandbook/13-vue-single-file-components\n4.0K vuehandbook/19-vue-components-communication\n84K vuehandbook/20-vue\n48K vuehandbook/21-bonus-vue-router\n1.2K vuehandbook/22-vue-practical-production\n128K vuehandbook/23-vue-first-app\n8.0K vuehandbook/96-vue-components\n536K vuehandbook/04-vue-devtools\nflavio@flavio-laptop:~$
```

Adding the `-a` option will print the size of each file in the directories, too:

```
flavio@flavio-laptop:~$ du -ah vuehandbook\n4.0K vuehandbook/12-vue-watchers/index.md\n4.0K vuehandbook/13-vue-watchers\n4.0K vuehandbook/07-vue-single-file-components/index.md\n4.0K vuehandbook/19-vue-components-communication/index.md\n4.0K vuehandbook/19-vue-components-communication\n36K vuehandbook/20-vue-router/vue-store.png\n36K vuehandbook/20-vue-router/vue-store\n36K vuehandbook/20-vue-router/Codesandbox.png\n84K vuehandbook/22-vue-router/bonus\n128K vuehandbook/23-vue-first-app\n32K vuehandbook/21-bonus-vue-router/banner.jpg\n16K vuehandbook/21-bonus-vue-router/index.md\nflavio@flavio-laptop:~$
```

A handy thing is to sort the directories by size:

```
du -h <directory> | sort -nr
```

and then piping to `head` to only get the first 10 results:

```
flavio@flavio-laptop:~$ du -h -b vuehandbook | sort -nr | head\n932K vuehandbook/05-vue-vcode\n636K vuehandbook/41/vt/objectives/75\n544K vuehandbook/04-vue-devtools\n536K vuehandbook/04-vue-devtools\n128K vuehandbook/02-vue-first-app\n88K vuehandbook/04-vue-devtools/pack\n88K vuehandbook/20-vue\n84K vuehandbook/20-vue\n76K vuehandbook/20-vue-store/67\n64K vuehandbook/20-vue-store/41\nflavio@flavio-laptop:~$
```

69

70

71

df

The `df` command is used to get disk usage information.

Its basic form will print information about the volumes mounted:

Using the `-h` option (`df -h`) will show those values in a human-readable format:

You can also specify a file or directory name to get information about the specific volume it lives on:

basename

Suppose you have a path to a file, for example `/Users/flavio/test.txt`.

Running

```
basename /Users/flavio/test.txt
```

will return the `test.txt` string:

If you run `basename` on a path string that points to a directory, you will get the last segment of the path. In this example, `/Users/flavio` is a directory:

72

73

dirname

Suppose you have a path to a file, for example `/Users/flavio/test.txt`.

Running

```
dirname /Users/flavio/test.txt
```

will return the `/Users/flavio` string:

74

You can search for a specific process combining `grep` with a pipe, like this:

```
ps axww | grep "Visual Studio Code"
```

The columns returned by `ps` represent some key information.

The first information is `PID`, the process ID. This is key when you want to reference this process in another command, for example to kill it.

Then we have `TT` that tells us the terminal id used.

Then `STAT` tells us the state of the process:

I a process that is idle (sleeping for longer than about 20 seconds) R a runnable process S a process that is sleeping for less than about 20 seconds T a stopped process U a process in uninterruptible wait Z a dead process (a zombie)

If you have more than one letter, the second represents further information, which can be very technical.

ps

Your computer is running, at all times, tons of different processes.

You can inspect them all using the `ps` command:

This is the list of user-initiated processes currently running in the current session.

Here I have a few `fish` shell instances, mostly opened by VS Code inside the editor, and an instances of Hugo running the development preview of a site.

Those are just the commands assigned to the current user. To list all processes we need to pass some options to `ps`.

The most common I use is `ps ax`:

The `a` option is used to also list other users processes, not just our own. `x` shows processes not linked to any terminal (not initiated by users through a terminal).

As you can see, the longer commands are cut. Use the command `ps axww` to continue the command listing on a new line instead of cutting it:

We need to specify `w` 2 times to apply this setting, it's not a typo.

75

76

77

It's common to have `+` which indicates the process is in the foreground in its terminal. `s` means the process is a [session leader](#).

`TIME` tells us how long the process has been running.

top

A quick guide to the `top` command, used to list the processes running in real time

The `top` command is used to display dynamic real-time information about running processes in the system.

It's really handy to understand what is going on.

Its usage is simple, you just type `top`, and the terminal will be fully immersed in this new view:

PID	COMMAND	%CPU	%MEM	TIME	%NICE	%IRQ	%MIGR	%RPS	%CPRS	%PSP	
229	NodeServer	16.1	13.11	1:18	4	6687.	84%	1218k	173k	229	
1133	nodecg-worker	0.0	0.0	0:00	0	4	100%	0	0	1133	
82651	top	7.4	0.8	0:01.48	37.1	0	27.	8352k	80	82651	
383	nodecg-worker	0.0	0.0	0:00.11	0	0	100%	0	0	383	
2152	Passenger	7	0.8	0:01.48	8	2	427.	159k	3	2152	
9	Kernel_Lock	3.9	0.6	0:15.45	262/12	0	0	510k	80	9	
82652	nodecg-worker	0.0	0.0	0:00.00	0	0	100%	0	0	82652	
377	Bear	2.1	41.88	1:12	5	1768-	57%	13k	468k	377	
4033	nodecg-worker	0.0	0.0	0:00.00	0	0	100%	0	0	4033	
347	Cloud	2.0	18.18	0:16	5	785-	13%	304k	12k	347	
549	Node	1.5	45.37	17.22	5	568-	79%	80k	240k	549	
548	Bluetooth	1.5	44.49	19.58	3	1	884	1%	60k	548	
1134	Recycleagle	1.3	87.32	1:23	3	1	218-	25%	80	434	1134

The process is long-running. To quit, you can type the `q` letter or `ctrl-C`.

There's a lot of information being given to us: the number of processes, how many are running or sleeping, the system load, the CPU usage, and a lot more.

Below, the list of processes taking the most memory and CPU is constantly updated.

By default, as you can see from the `%CPU` column highlighted, they are sorted by the CPU used.

You can add a flag to sort processes by memory utilized:

```
top -o mem
```

kill

Linux processes can receive [signals](#) and react to them.

That's one way we can interact with running programs.

The `kill` program can send a variety of signals to a program.

It's not just used to terminate a program, like the name would suggest, but that's its main job.

We use it in this way:

```
kill <PID>
```

By default, this sends the `TERM` signal to the process id specified.

We can use flags to send other signals, including:

```
kill -HUP <PID>
kill -INT <PID>
kill -KILL <PID>
kill -TERM <PID>
kill -CONT <PID>
kill -STOP <PID>
```

`HUP` means [hang up](#). It's sent automatically when a terminal window that started a process is closed before terminating the process.

`INT` means [interrupt](#), and it sends the same signal used when we press `ctrl-C` in the terminal, which usually terminates the process.

`KILL` is not sent to the process, but to the operating system kernel, which immediately stops and terminates the process.

`TERM` means [terminate](#). The process will receive it and terminate itself. It's the default signal sent by `kill`.

`CONT` means [continue](#). It can be used to resume a stopped process.

`STOP` is not sent to the process, but to the operating system kernel, which immediately stops (but does not terminate) the process.

You might see numbers used instead, like `kill -1 <PID>`. In this case,

```
1 corresponds to HUP . 2 corresponds to INT . 9
corresponds to KILL . 15 corresponds to TERM .
18 corresponds to CONT . 15 corresponds to
STOP .
```

killall

Similar to the `kill` command, `killall` instead of sending a signal to a specific process id will send the signal to multiple processes at once.

This is the syntax:

```
killall <name>
```

where `name` is the name of a program. For example you can have multiple instances of the `top` program running, and `killall top` will terminate them all.

You can specify the signal, like with `kill` (and check the `kill` tutorial to read more about the specific kinds of signals we can send), for example:

```
killall -HUP top
```

jobs

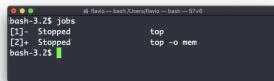
When we run a command in Linux / macOS, we can set it to run in the background using the `&` symbol after the command. For example we can run `top` in the background:

```
top &
```

This is very handy for long-running programs.

We can get back to that program using the `fg` command. This works fine if we just have one job in the background, otherwise we need to use the job number: `fg 1`, `fg 2` and so on. To get the job number, we use the `jobs` command.

Say we run `top &` and then `top -o mem &`, so we have 2 top instances running. `jobs` will tell us this:



```
bat@MacBook-Pro ~ % bat@MacBook-Pro ~ % top &
bat@MacBook-Pro ~ % top -o mem &
bat@MacBook-Pro ~ % jobs
[1]+  Stopped                  top (wd:-)
[2]+  Stopped                  top -o mem (wd:-)
```

Now we can switch back to one of those using `fg <jobid>`. To stop the program again we can hit `cmd-Z`.

Running `jobs -l` will also print the process id of each job.

bg

When a command is running you can suspend it using `ctrl-Z`.

The command will immediately stop, and you get back to the shell terminal.

You can resume the execution of the command in the background, so it will keep running but it will not prevent you from doing other work in the terminal.

In this example I have 2 commands stopped:



```
bat@MacBook-Pro ~ % bat@MacBook-Pro ~ % bg 1
bat@MacBook-Pro ~ % bg
bat@MacBook-Pro ~ % jobs
[1]+  Stopped                  top (wd:-)
[2]+  Stopped                  top -o mem (wd:-)
```

I can run `bg 1` to resume in the background the execution of the job #1.

I could have also said `bg` without any option, as the default is to pick the job #1 in the list.

fg

When a command is running in the background, because you started it with `&` at the end (example: `top &`) or because you put it in the background with the `bg` command, you can put it to the foreground using `fg`.

Running

```
fg
```

will resume to the foreground the last job that was suspended.

You can also specify which job you want to resume to the foreground passing the job number, which you can get using the `jobs` command.



```
bat@MacBook-Pro ~ % bat@MacBook-Pro ~ % fg 2
bat@MacBook-Pro ~ % fg
bat@MacBook-Pro ~ % jobs
[1]+  Stopped                  top (wd:-)
[2]+  Stopped                  top -o mem (wd:-)
```

Running `fg 2` will resume job #2:

84

85

86

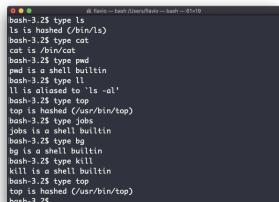
type

A command can be one of those 4 types:

- an executable
- a shell built-in program
- a shell function
- an alias

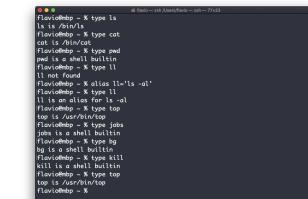
The `type` command can help figure out this, in case we want to know or we're just curious. It will tell you how the command will be interpreted.

The output will depend on the shell used. This is Bash:



```
bat@MacBook-Pro ~ % bat@MacBook-Pro ~ % type ls
ls is hashed (/bin/ls)
bat@MacBook-Pro ~ % type cat
cat is a shell function
bat@MacBook-Pro ~ % type pwd
pwd is a shell builtin
bat@MacBook-Pro ~ % type top
top is a shell function
bat@MacBook-Pro ~ % type kill
kill is a shell builtin
bat@MacBook-Pro ~ % type top
top is /usr/bin/top
bat@MacBook-Pro ~ %
```

This is Zsh:



```
bat@MacBook-Pro ~ % bat@MacBook-Pro ~ % type ls
ls is a function with definition
# defined in /usr/local/Cellar/fish/3.1.0/share/fish/functions/ls.fish @ line 13
function ls {
    set -l opt -E
    set -o extendedglob
    set -o setopt -F
    command ls $opt $arg
}
end
bat@MacBook-Pro ~ % type cat
cat is a function
# defined in /usr/local/Cellar/fish/3.1.0/share/fish/functions/cat.fish @ line 11
function cat {
    set -l opt -E
    set -o extendedglob
    set -o setopt -F
    command cat $opt $arg
}
end
bat@MacBook-Pro ~ % type pwd
pwd is a shell builtin
bat@MacBook-Pro ~ % type top
top is /usr/bin/top
bat@MacBook-Pro ~ % type kill
kill is a shell builtin
bat@MacBook-Pro ~ % type top
top is /usr/bin/top
bat@MacBook-Pro ~ %
```

This is Fish:



```
bat@MacBook-Pro ~ % bat@MacBook-Pro ~ % type ls
ls is a function with definition
# defined in /usr/local/Cellar/fish/3.1.0/share/fish/functions/ls.fish @ line 13
function ls {
    set -l opt -E
    set -o extendedglob
    set -o setopt -F
    command ls $opt $arg
}
end
bat@MacBook-Pro ~ % type cat
cat is a function
# defined in /usr/local/Cellar/fish/3.1.0/share/fish/functions/cat.fish @ line 11
function cat {
    set -l opt -E
    set -o extendedglob
    set -o setopt -F
    command cat $opt $arg
}
end
bat@MacBook-Pro ~ % type pwd
pwd is a shell builtin
bat@MacBook-Pro ~ % type top
top is /usr/bin/top
bat@MacBook-Pro ~ % type kill
kill is a function with definition
# defined in /usr/local/Cellar/fish/3.1.0/share/fish/functions/kill.fish @ line 4
function kill {
    local pid=$1
    shift
    command kill $pid ${@:1}
}
end
bat@MacBook-Pro ~ % type top
top is /usr/bin/top
bat@MacBook-Pro ~ %
```

One of the most interesting things here is that for aliases it will tell you what is aliasing to. You can see the `ll` alias, in the case of Bash and Zsh, but Fish provides it by default, so it will tell you it's a built-in shell function.

which

Suppose you have a command you can execute, because it's in the shell path, but you want to know where it is located.

You can do so using `which`. The command will return the path to the command specified:



```
+ ~ which ls  
/bin/ls  
+ ~ which docker  
/usr/local/bin/docker  
+ ~
```

`which` will only work for executables stored on disk, not aliases or built-in shell functions.

90

xargs

The `xargs` command is used in a UNIX shell to convert input from standard input into arguments to a command.

In other words, through the use of `xargs` the output of a command is used as the input of another command.

Here's the syntax you will use:

```
command1 | xargs command2
```

We use a pipe (`|`) to pass the output to `xargs`. That will take care of running the `command2` command, using the output of `command1` as its argument(s).

Let's do a simple example. You want to remove some specific files from a directory. Those files are listed inside a text file.

We have 3 files: `file1`, `file2`, `file3`.

In `todelete.txt` we have a list of files we want to delete, in this example `file1` and `file3`:



```
+ testing ls  
file1 file2 file3 todelete.txt  
+ testing cat todelete.txt  
file1  
file3  
+ testing ls  
file2 todelete.txt  
+ testing
```

91

We will channel the output of `cat todelete.txt` to the `rm` command, through `xargs`.

In this way:

```
cat todelete.txt | xargs rm
```

That's the result, the files we listed are now deleted:



```
+ testing ls  
file1 file2 file3 todelete.txt  
+ testing cat todelete.txt | xargs rm  
file1  
file3  
+ testing ls  
file2 todelete.txt  
+ testing
```

The way it works is that `xargs` will run `rm` 2 times, one for each line returned by `cat`.

This is the simplest usage of `xargs`. There are several options we can use.

One of the most useful in my opinion, especially when starting to learn `xargs`, is `-p`. Using this option will make `xargs` print a confirmation prompt with the action it's going to take:



```
+ testing ls  
file1 file2 file3 todelete.txt  
+ testing cat todelete.txt | xargs -p rm  
rm file1 file3?...|
```

nohup

Sometimes you have to run a long-lived process on a remote machine, and then you need to disconnect.

Or you simply want to prevent the command to be halted if there's any network issue between you and the server.

The way to make a command run even after you log out or close the session to a server is to use the `nohup` command.

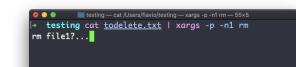
Use `nohup <command>` to let the process continue working even after you log out.

93

94

92

The `-n` option lets you tell `xargs` to perform one iteration at a time, so you can individually confirm them with `-p`. Here we tell `xargs` to perform one iteration at a time with `-n1`:



```
+ testing cat todelete.txt | xargs -p -n1 rm  
rm file1?...|
```

The `-I` option is another widely used one. It allows you to get the output into a placeholder, and then you can do various things.

One of them is to run multiple commands:

```
command1 | xargs -I % /bin/bash -c 'command2 %; com
```



```
+ testing cat todelete.txt | xargs -p -I % sh -c 'ls %; rm %'  
sh -c 'ls file1; rm file1?...|
```

You can swap the `%` symbol I used above with anything else, it's a variable

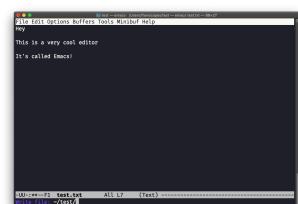
95

macOS users, stop a second now. If you are on Linux there are no problems, but macOS does not ship applications using GPLv3, and every built-in UNIX command that has been updated to GPLv3 has not been updated. While there is a little problem with the commands I listed up to now, in this case using an emacs version from 2007 is not exactly the same as using a version with 12 years of improvements and change. This is not a problem with Vim, which is up to date. To fix this, run `brew install emacs` and running `emacs` will use the new version from Homebrew (make sure you have Homebrew installed)

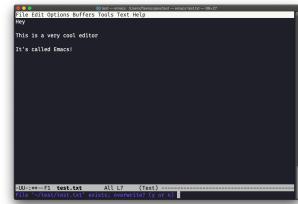
You can also edit an existing file calling `emacs <filename>`:



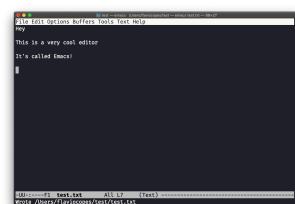
You can start editing and once you are done, press `ctrl-x` followed by `ctrl-w`. You confirm the folder:



and Emacs tell you the file exists, asking you if it should overwrite it:



Answer `y`, and you get a confirmation of success:



You can exit Emacs pressing `ctrl-x` followed by `ctrl-c`. Or `ctrl-x` followed by `c` (keep `ctrl` pressed).

There is a lot to know about Emacs. More than I am able to write in this little introduction. I encourage you to open Emacs and press `ctrl-h r` to open the built-in manual and `ctrl-h t` to open the official tutorial.

102

103

104

nano

`nano` is a beginner friendly editor.

Run it using `nano <filename>`.

You can directly type characters into the file without worrying about modes.

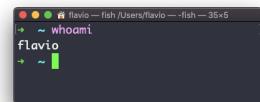
You can quit without editing using `ctrl-x`. If you edited the file buffer, the editor will ask you for confirmation and you can save the edits, or discard them. The help at the bottom shows you the keyboard commands that let you work with the file:



`pico` is more or less the same, although `nano` is the GNU version of `pico` which at some point in history was not open source and the `nano` clone was made to satisfy the GNU operating system license requirements.

whoami

Type `whoami` to print the user name currently logged in to the terminal session:

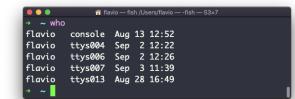


Note: this is different from the `who am i` command, which prints more information

who

The `who` command displays the users logged in to the system.

Unless you're using a server multiple people have access to, chances are you will be the only user logged in, multiple times:



Why multiple times? Because each shell opened will count as an access.

You can see the name of the terminal used, and the time/day the session was started.

The `-aft` flags will tell `who` to display more information, including the idle time and the process ID of the terminal:

105

106

107

```

USER  LINE    WHEN      IDLE   PID COMMENT
flavio  console Aug 13 12:25:30 old    167
flavio  ttys000 Aug 13 12:25:30 old    70258 term=0 exit=0
flavio  ttys000 Sep  2 12:26:00 old    71201 term=0 exit=0
flavio  ttys000 Sep  2 12:26:00:41 71675
flavio  ttys007 Sep  3 11:39:00:01 82856
flavio  ttys008 Aug 21 14:31:04:15 72016 term=0 exit=0
flavio  ttys008 Aug 21 14:31:04:15 72214 term=0 exit=0
flavio  ttys018 Aug 23 16:30:32 38649 term=0 exit=0
flavio  ttys018 Aug 23 16:30:32 38855 term=0 exit=0
flavio  ttys018 Aug 23 16:49:01:33 5086 run-level 3

```

The special `who am i` command will list the current terminal session details:

```

flavio  ttys004 Sep  2 12:22

```

```

USER  LINE    WHEN      IDLE   PID COMMENT
flavio  console Aug 13 12:25:30 old    71365

```

SU

While you're logged in to the terminal shell with one user, you might have the need to switch to another user.

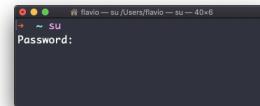
For example you're logged in as root to perform some maintenance, but then you want to switch to a user account.

You can do so with the `su` command:

```
su <username>
```

For example: `su flavio`.

If you're logged in as a user, running `su` without anything else will prompt to enter the `root` user password, as that's the default behavior.



`su` will start a new shell as another user.

When you're done, typing `exit` in the shell will close that shell, and will return back to the current user's shell.

sudo

`sudo` is commonly used to run a command as root.

You must be enabled to use `sudo`, and once you do, you can run commands as root by entering your user's password (not the root user password).

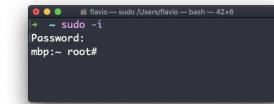
The permissions are highly configurable, which is great especially in a multi-user server environment, and some users can be granted access to running specific commands through `sudo`.

For example you can edit a system configuration file:

```
sudo nano /etc/hosts
```

which would otherwise fail to save since you don't have the permissions for it.

You can run `sudo -i` to start a shell as root:



You can use `sudo` to run commands as any user. `root` is the default, but use the `-u` option to specify another user:

108

109

110

```
sudo -u flavio ls /Users/flavio
```

passwd

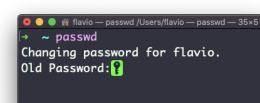
Users in Linux have a password assigned. You can change the password using the `passwd` command.

There are two situations here.

The first is when you want to change your password. In this case you type:

```
passwd
```

and an interactive prompt will ask you for the old password, then it will ask you for the new one:



When you're `root` (or have superuser privileges) you can set the username of which you want to change the password:

```
passwd <username> <new password>
```

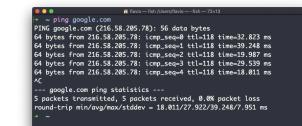
In this case you don't need to enter the old one.

ping

The `ping` command pings a specific network host, on the local network or on the Internet.

You use it with the syntax `ping <host>` where `<host>` could be a domain name, or an IP address.

Here's an example pinging `google.com`:



The commands sends a request to the server, and the server returns a response.

`ping` keep sending the request every second, by default, and will keep running until you stop it with `ctrl-C`, unless you pass the number of times you want to try with the `-c` option: `ping -c 2 google.com`.

Once `ping` is stopped, it will print some statistics about the results: the percentage of packages lost, and statistics about the network performance.

As you can see the screen prints the host IP address, and the time that it took to get the response back.

111

112

113

Not all servers support pinging, in case the requests times out:

```
PING flaviocopes.com (167.99.137.12): 56 data bytes
Request timeout for icmp_seq 0
Request timeout for icmp_seq 1
Request timeout for icmp_seq 2
Request timeout for icmp_seq 3
...
-- flaviocopes.com ping statistics --
5 packets transmitted, 0 packets received, 100% packet loss
```

Sometimes this is done on purpose, to "hide" the server, or just to reduce the load. The ping packets can also be filtered by firewalls.

ping works using the **ICMP protocol** (*Internet Control Message Protocol*), a network layer protocol just like TCP or UDP.

The request sends a packet to the server with the `ECHO_REQUEST` message, and the server returns a `ECHO_REPLY` message. I won't go into details, but this is the basic concept.

Pinging a host is useful to know if the host is reachable (supposing it implements ping), and how distant it is in terms of how long it takes to get back to you. Usually the nearest server is geographically, the less time it will take to return back to you, for simple physical laws that cause a longer distance to introduce more delay in the cables.

traceroute

When you try to reach a host on the Internet, you go through your home router, then you reach your ISP network, which in turn goes through its own upstream network router, and so on, until you finally reach the host.

Have you ever wanted to know what are the steps that your packets go through to do that?

The traceroute command is made for this.

You invoke

```
traceroute <host>
```

and it will (slowly) gather all the information while the packet travels.

In this example I tried reaching for my blog with traceroute flaviocopes.com :

```
traceroute to flaviocopes.com (167.99.137.12)
traceroute: Warning: flaviocopes.com has multiple addresses; using 167.99.138.123
traceroute to flaviocopes.com (167.99.137.123), 30 hops max, 52 byte packets
1 192.168.1.1 (192.168.1.1) 6.404 ms
2 192.168.1.10 (192.168.1.10) 10.697 ms 20.080 ms 21.329 ms
3 192.168.1.10 (192.168.1.10) 10.697 ms 20.080 ms 21.329 ms
4 192.168.1.10 (192.168.1.10) 10.697 ms 20.080 ms 21.329 ms
5 192.168.1.10 (192.168.1.10) 10.697 ms 20.080 ms 21.329 ms
6 192.168.1.10 (192.168.1.10) 10.697 ms 20.080 ms 21.329 ms
7 fwe0-1-link.telia.net (82.15.15.172) 28.671 ms 39.589 ms *
8 fwe0-1-link.telia.net (82.15.15.172) 28.671 ms 39.589 ms *
9 digitaleocon-ic-328177-fwe0-1c.telia.net (80.239.128.23) 172.565 ms 173.744 ms
digitaleocon-ic-328177-fwe0-1c.telia.net (80.239.128.23) 156.124 ms
10 *
11 *
12 *
13 162.99.138.123 (162.99.138.123) 287.791 ms 12 32.786 ms 12 23.561 ms 12
14 162.99.138.123 (162.99.138.123) 35.460 ms 12
```

Not every router travelled returns us information. In this case, traceroute prints * * *. Otherwise, we can see the hostname, the IP address, and some performance indicator.

For every router we can see 3 samples, which means traceroute tries by default 3 times to get you a good indication of the time needed to reach it. This is why it takes this long to execute traceroute compared to simply doing a ping to that host.

You can customize this number with the `-q` option:

```
traceroute -q 1 flaviocopes.com
```

```
traceroute to flaviocopes.com (167.99.137.12)
traceroute: warning: flaviocopes.com has multiple addresses; using 167.99.137.12
traceroute to flaviocopes.com (167.99.137.12), 30 hops max, 52 byte packets
1 192.168.1.1 (192.168.1.1) 6.401 ms
2 192.168.1.10 (192.168.1.10) 10.693 ms 17.438 ms
3 192.168.1.10 (192.168.1.10) 10.693 ms 17.438 ms
4 192.168.1.10 (192.168.1.10) 10.693 ms 17.438 ms
5 192.168.1.10 (192.168.1.10) 10.693 ms 17.438 ms
6 192.168.1.10 (192.168.1.10) 10.693 ms 17.438 ms
7 fwe0-1-link.telia.net (82.15.15.172) 28.677 ms 29.512 ms
8 fwe0-1-link.telia.net (82.15.15.172) 28.677 ms 29.512 ms
9 digitaleocon-ic-328177-fwe0-1c.telia.net (80.239.128.23) 137.479 ms
10 *
11 *
12 *
13 162.99.137.12 (162.99.137.12) 35.460 ms 12
```

114

115

116

clear

Type clear to clear all the previous commands that were ran in the current terminal.

The screen will clear and you will just see the prompt at the top:

```
flavi -- Flavi's Macbook Pro: ~ dh -- 66x10
```

Note: this command has a handy shortcut: `ctrl-L`

Once you do that, you will lose access to scrolling to see the output of the previous commands entered.

So you might want to use clear -x instead, which still clears the screen, but lets you go back to see the previous work by scrolling up.

history

Every time we run a command, that's memorized in the history.

You can display all the history using:

```
history
```

This shows the history with numbers:

```
111 passed
112 ls -l
113 ls -l
114 llcd mc
115 llcd mc
116 rm -v test.txt
117 echo test >> test.txt
118 wc test.txt
119 wc -l test.txt
120 ls -al | wc -l
121 ls -al | wc -l
122 rm -v test.txt
123 open .
124 !ls -al test.txt
125 !llcd history
126 history
127 history
128 history
```

You can use the syntax `!<command number>` to repeat a command stored in the history, in the above example typing `!121` will repeat the `ls -al | wc -l` command.

Typically the last 500 commands are stored in the history.

You can combine this with grep to find a command you ran:

```
history | grep docker
```

```
bash-3.2$ history | grep docker
7 git clone https://github.com/docker/getting-started.git
8 cd getting-started
9 docker container stop $(docker container ls)
10 docker container rm $(docker container ls)
11 docker container rm $(docker container ls -q)
12 history | grep docker
13 history | grep docker
14 history | grep docker
15 history | grep docker
16 history | grep docker
17 history | grep docker
18 history | grep docker
19 history | grep docker
20 history | grep docker
21 history | grep docker
22 history | grep docker
23 history | grep docker
24 history | grep docker
25 history | grep docker
26 history | grep docker
27 history | grep docker
28 history | grep docker
29 history | grep docker
30 history | grep docker
31 history | grep docker
32 history | grep docker
33 history | grep docker
34 history | grep docker
35 history | grep docker
36 history | grep docker
37 history | grep docker
38 history | grep docker
39 history | grep docker
40 history | grep docker
41 history | grep docker
42 history | grep docker
43 history | grep docker
44 history | grep docker
45 history | grep docker
46 history | grep docker
47 history | grep docker
48 history | grep docker
49 history | grep docker
50 history | grep docker
51 history | grep docker
52 history | grep docker
53 history | grep docker
54 history | grep docker
55 history | grep docker
56 history | grep docker
57 history | grep docker
58 history | grep docker
59 history | grep docker
60 history | grep docker
61 history | grep docker
62 history | grep docker
63 history | grep docker
64 history | grep docker
65 history | grep docker
66 history | grep docker
67 history | grep docker
68 history | grep docker
69 history | grep docker
70 history | grep docker
71 history | grep docker
72 history | grep docker
73 history | grep docker
74 history | grep docker
75 history | grep docker
76 history | grep docker
77 history | grep docker
78 history | grep docker
79 history | grep docker
80 history | grep docker
81 history | grep docker
82 history | grep docker
83 history | grep docker
84 history | grep docker
85 history | grep docker
86 history | grep docker
87 history | grep docker
88 history | grep docker
89 history | grep docker
90 history | grep docker
91 history | grep docker
92 history | grep docker
93 history | grep docker
94 history | grep docker
95 history | grep docker
96 history | grep docker
97 history | grep docker
98 history | grep docker
99 history | grep docker
100 history | grep docker
101 history | grep docker
102 history | grep docker
103 history | grep docker
104 history | grep docker
105 history | grep docker
106 history | grep docker
107 history | grep docker
108 history | grep docker
109 history | grep docker
110 history | grep docker
111 history | grep docker
112 history | grep docker
113 history | grep docker
114 history | grep docker
115 history | grep docker
116 history | grep docker
117 history | grep docker
118 history | grep docker
119 history | grep docker
120 history | grep docker
121 history | grep docker
122 history | grep docker
123 history | grep docker
124 history | grep docker
125 history | grep docker
126 history | grep docker
127 history | grep docker
128 history | grep docker
129 history | grep docker
130 history | grep docker
131 history | grep docker
132 history | grep docker
133 history | grep docker
134 history | grep docker
135 history | grep docker
136 history | grep docker
137 history | grep docker
138 history | grep docker
139 history | grep docker
140 history | grep docker
141 history | grep docker
142 history | grep docker
143 history | grep docker
144 history | grep docker
145 history | grep docker
146 history | grep docker
147 history | grep docker
148 history | grep docker
149 history | grep docker
150 history | grep docker
151 history | grep docker
152 history | grep docker
153 history | grep docker
154 history | grep docker
155 history | grep docker
156 history | grep docker
157 history | grep docker
158 history | grep docker
159 history | grep docker
160 history | grep docker
161 history | grep docker
162 history | grep docker
163 history | grep docker
164 history | grep docker
165 history | grep docker
166 history | grep docker
167 history | grep docker
168 history | grep docker
169 history | grep docker
170 history | grep docker
171 history | grep docker
172 history | grep docker
173 history | grep docker
174 history | grep docker
175 history | grep docker
176 history | grep docker
177 history | grep docker
178 history | grep docker
179 history | grep docker
180 history | grep docker
181 history | grep docker
182 history | grep docker
183 history | grep docker
184 history | grep docker
185 history | grep docker
186 history | grep docker
187 history | grep docker
188 history | grep docker
189 history | grep docker
190 history | grep docker
191 history | grep docker
192 history | grep docker
193 history | grep docker
194 history | grep docker
195 history | grep docker
196 history | grep docker
197 history | grep docker
198 history | grep docker
199 history | grep docker
200 history | grep docker
201 history | grep docker
202 history | grep docker
203 history | grep docker
204 history | grep docker
205 history | grep docker
206 history | grep docker
207 history | grep docker
208 history | grep docker
209 history | grep docker
210 history | grep docker
211 history | grep docker
212 history | grep docker
213 history | grep docker
214 history | grep docker
215 history | grep docker
216 history | grep docker
217 history | grep docker
218 history | grep docker
219 history | grep docker
220 history | grep docker
221 history | grep docker
222 history | grep docker
223 history | grep docker
224 history | grep docker
225 history | grep docker
226 history | grep docker
227 history | grep docker
228 history | grep docker
229 history | grep docker
230 history | grep docker
231 history | grep docker
232 history | grep docker
233 history | grep docker
234 history | grep docker
235 history | grep docker
236 history | grep docker
237 history | grep docker
238 history | grep docker
239 history | grep docker
240 history | grep docker
241 history | grep docker
242 history | grep docker
243 history | grep docker
244 history | grep docker
245 history | grep docker
246 history | grep docker
247 history | grep docker
248 history | grep docker
249 history | grep docker
250 history | grep docker
251 history | grep docker
252 history | grep docker
253 history | grep docker
254 history | grep docker
255 history | grep docker
256 history | grep docker
257 history | grep docker
258 history | grep docker
259 history | grep docker
260 history | grep docker
261 history | grep docker
262 history | grep docker
263 history | grep docker
264 history | grep docker
265 history | grep docker
266 history | grep docker
267 history | grep docker
268 history | grep docker
269 history | grep docker
270 history | grep docker
271 history | grep docker
272 history | grep docker
273 history | grep docker
274 history | grep docker
275 history | grep docker
276 history | grep docker
277 history | grep docker
278 history | grep docker
279 history | grep docker
280 history | grep docker
281 history | grep docker
282 history | grep docker
283 history | grep docker
284 history | grep docker
285 history | grep docker
286 history | grep docker
287 history | grep docker
288 history | grep docker
289 history | grep docker
290 history | grep docker
291 history | grep docker
292 history | grep docker
293 history | grep docker
294 history | grep docker
295 history | grep docker
296 history | grep docker
297 history | grep docker
298 history | grep docker
299 history | grep docker
300 history | grep docker
301 history | grep docker
302 history | grep docker
303 history | grep docker
304 history | grep docker
305 history | grep docker
306 history | grep docker
307 history | grep docker
308 history | grep docker
309 history | grep docker
310 history | grep docker
311 history | grep docker
312 history | grep docker
313 history | grep docker
314 history | grep docker
315 history | grep docker
316 history | grep docker
317 history | grep docker
318 history | grep docker
319 history | grep docker
320 history | grep docker
321 history | grep docker
322 history | grep docker
323 history | grep docker
324 history | grep docker
325 history | grep docker
326 history | grep docker
327 history | grep docker
328 history | grep docker
329 history | grep docker
330 history | grep docker
331 history | grep docker
332 history | grep docker
333 history | grep docker
334 history | grep docker
335 history | grep docker
336 history | grep docker
337 history | grep docker
338 history | grep docker
339 history | grep docker
340 history | grep docker
341 history | grep docker
342 history | grep docker
343 history | grep docker
344 history | grep docker
345 history | grep docker
346 history | grep docker
347 history | grep docker
348 history | grep docker
349 history | grep docker
350 history | grep docker
351 history | grep docker
352 history | grep docker
353 history | grep docker
354 history | grep docker
355 history | grep docker
356 history | grep docker
357 history | grep docker
358 history | grep docker
359 history | grep docker
360 history | grep docker
361 history | grep docker
362 history | grep docker
363 history | grep docker
364 history | grep docker
365 history | grep docker
366 history | grep docker
367 history | grep docker
368 history | grep docker
369 history | grep docker
370 history | grep docker
371 history | grep docker
372 history | grep docker
373 history | grep docker
374 history | grep docker
375 history | grep docker
376 history | grep docker
377 history | grep docker
378 history | grep docker
379 history | grep docker
380 history | grep docker
381 history | grep docker
382 history | grep docker
383 history | grep docker
384 history | grep docker
385 history | grep docker
386 history | grep docker
387 history | grep docker
388 history | grep docker
389 history | grep docker
390 history | grep docker
391 history | grep docker
392 history | grep docker
393 history | grep docker
394 history | grep docker
395 history | grep docker
396 history | grep docker
397 history | grep docker
398 history | grep docker
399 history | grep docker
400 history | grep docker
401 history | grep docker
402 history | grep docker
403 history | grep docker
404 history | grep docker
405 history | grep docker
406 history | grep docker
407 history | grep docker
408 history | grep docker
409 history | grep docker
410 history | grep docker
411 history | grep docker
412 history | grep docker
413 history | grep docker
414 history | grep docker
415 history | grep docker
416 history | grep docker
417 history | grep docker
418 history | grep docker
419 history | grep docker
420 history | grep docker
421 history | grep docker
422 history | grep docker
423 history | grep docker
424 history | grep docker
425 history | grep docker
426 history | grep docker
427 history | grep docker
428 history | grep docker
429 history | grep docker
430 history | grep docker
431 history | grep docker
432 history | grep docker
433 history | grep docker
434 history | grep docker
435 history | grep docker
436 history | grep docker
437 history | grep docker
438 history | grep docker
439 history | grep docker
440 history | grep docker
441 history | grep docker
442 history | grep docker
443 history | grep docker
444 history | grep docker
445 history | grep docker
446 history | grep docker
447 history | grep docker
448 history | grep docker
449 history | grep docker
450 history | grep docker
451 history | grep docker
452 history | grep docker
453 history | grep docker
454 history | grep docker
455 history | grep docker
456 history | grep docker
457 history | grep docker
458 history | grep docker
459 history | grep docker
460 history | grep docker
461 history | grep docker
462 history | grep docker
463 history | grep docker
464 history | grep docker
465 history | grep docker
466 history | grep docker
467 history | grep docker
468 history | grep docker
469 history | grep docker
470 history | grep docker
471 history | grep docker
472 history | grep docker
473 history | grep docker
474 history | grep docker
475 history | grep docker
476 history | grep docker
477 history | grep docker
478 history | grep docker
479 history | grep docker
480 history | grep docker
481 history | grep docker
482 history | grep docker
483 history | grep docker
484 history | grep docker
485 history | grep docker
486 history | grep docker
487 history | grep docker
488 history | grep docker
489 history | grep docker
490 history | grep docker
491 history | grep docker
492 history | grep docker
493 history | grep docker
494 history | grep docker
495 history | grep docker
496 history | grep docker
497 history | grep docker
498 history | grep docker
499 history | grep docker
500 history | grep docker
```

To clear the history, run `history -c`

114

115

116

Note: this command has a handy shortcut: `ctrl-L`

Once you do that, you will lose access to scrolling to see the output of the previous commands entered.

So you might want to use `clear -x` instead, which still clears the screen, but lets you go back to see the previous work by scrolling up.

You can use the syntax `!<command number>` to repeat a command stored in the history, in the above example typing `!121` will repeat the `ls -al | wc -l` command.

Typically the last 500 commands are stored in the history.

You can combine this with grep to find a command you ran:

```
history | grep docker
```

117

118

119

export

The `export` command is used to export variables to child processes.

What does this mean?

Suppose you have a variable `TEST` defined in this way:

```
TEST="test"
```

You can print its value using `echo $TEST`:

```
flavio@flavio-MacBook-Pro:~$ echo $TEST
test
bash:3.25
```

But if you try defining a Bash script in a file `script.sh` with the above command:

```
flavio@flavio-MacBook-Pro:~/Desktop$ nano script.sh
#!/bin/bash
echo $TEST

[ Read 2 lines ]
M Get Help W WriteOut R Read File P Prev PgUp K Cut Text C Cur Pos
X Exit J Justify H Where I Up Next PgDn Uncat T To Spell
```

Then you set `chmod u+x script.sh` and you execute this script with `./script.sh`, the `echo $TEST` line will print nothing!

This is because in Bash the `TEST` variable was defined local to the shell. When executing a shell script or another command, a subshell is launched to execute it, which does not contain the current shell local variables.

To make the variable available there we need to define `TEST` not in this way:

```
TEST="test"
```

but in this way:

```
export TEST="test"
```

Try that, and running `./script.sh` now should print "test":

```
flavio@flavio-MacBook-Pro:~/Desktop$ export TEST="test"
flavio@flavio-MacBook-Pro:~/Desktop$ ./script.sh
test
bash:3.25
```

Sometimes you need to append something to a variable. It's often done with the `PATH` variable. You use this syntax:

```
export PATH=$PATH:/new/path
```

120

121

crontab

Cron jobs are jobs that are scheduled to run at specific intervals. You might have a command perform something every hour, or every day, or every 2 weeks. Or on weekends. They are very powerful, especially on servers to perform maintenance and automations.

The `crontab` command is the entry point to work with cron jobs.

The first thing you can do is to explore which cron jobs are defined by me:

```
crontab -l
```

You might have none, like me:

```
flavio@flavio-MacBook-Pro:~$ crontab -l
crontab: no crontab for flavio
: ~
```

Run

```
crontab -e
```

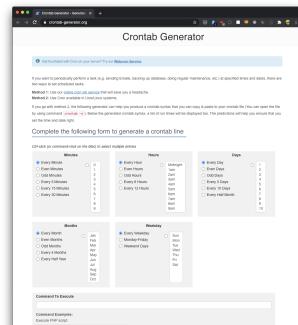
to edit the cron jobs, and add new ones.

By default this opens with the default editor, which is usually `vim`. I like `nano` more, you can use this line to use a different editor:

```
EDITOR=nano crontab -e
```

Now you can add one line for each cron job.

The syntax to define cron jobs is kind of scary. This is why I usually use a website to help me generate it without errors: <https://crontab-generator.org/>



You pick a time interval for the cron job, and you type the command to execute.

It's common to use `export` when you create new variables in this way, but also when you create variables in the `.bash_profile` or `.bashrc` configuration files with Bash, or in `.zshenv` with Zsh.

To remove a variable, use the `-n` option:

```
export -n TEST
```

Calling `export` without any option will list all the exported variables.

122

I chose to run a script located in `/Users/flavio/test.sh` every 12 hours. This is the crontab line I need to run:

```
* *12 * * * /Users/flavio/test.sh >/dev/null 2>1
```

I run `crontab -e`:

```
EDITOR=nano crontab -e
```

and I add that line, then I press `ctrl-X` and press `y` to save.

If all goes well, the cron job is set up:

```
flavio@flavio-MacBook-Pro:~$ bash:3.25 EDITOR=nano crontab -e
crontab: no crontab for flavio - using an empty one
crontab: installing new crontab
bash:3.25
bash:3.25
```

Once this is done, you can see the list of active cron jobs by running:

```
crontab -l
```

```
flavio@flavio-MacBook-Pro:~$ bash:3.25 crontab -l
* *12 * * * /Users/flavio/test.sh >/dev/null 2>1
bash:3.25
```

123

124

125

You can remove a cron job running `crontab -e` again, removing the line and exiting the editor:

```
GNU nano 2.0.6  File: /tmp/crontab_30t2K2jnTC Modified
Save modified buffer (ANSWERING "No" WILL DESTROY CHANGES)? Y
Y Yes
N No
F Cancel
```

```
bash-3.2$ EDITOR=nano crontab -e
crontab: installing new crontab
bash-3.2$ crontab -l
bash-3.2$
```

uname

Calling `uname` without any options will return the Operating System codename:

```
flavio@flavio: ~
flavio@flavio: ~$ uname
Darwin
flavio@flavio: ~
```

The `m` option shows the hardware name (`x86_64` in this example) and the `p` option prints the processor architecture name (`i386` in this example):

```
flavio@flavio: ~
flavio@flavio: ~$ uname -mp
x86_64 i386
flavio@flavio: ~
```

The `s` option prints the Operating System name. `r` prints the release, `v` prints the version:

```
flavio@flavio: ~
flavio@flavio: ~$ uname -srsv
Darwin 19.6.0 Darwin Kernel Version 19.6.0:
Thu Jun 18 20:49:00 PDT 2020; root:xnu-6153.141.1~1/RELEASE_X86_64
flavio@flavio: ~
```

The `n` option prints the node network name:

```
flavio@flavio: ~
flavio@flavio: ~$ uname -n
mbp.local
flavio@flavio: ~
```

The `a` option prints all the information available:

```
flavio@flavio: ~
flavio@flavio: ~$ uname -a
Darwin mbp.local 19.6.0 Darwin Kernel Version 19.6.0:
Thu Jun 18 20:49:00 PDT 2020; root:xnu-6153.141.1~1/RELEASE_X86_64 x86_64
flavio@flavio: ~
```

126

127

128

On macOS you can also use the `sw_vers` command to print more information about the macOS Operating System. Note that this differs from the Darwin (the Kernel) version, which above is `19.6.0`.

Darwin is the name of the kernel of macOS. The kernel is the "core" of the Operating System, while the Operating System as a whole is called macOS. In Linux, Linux is the kernel, GNU/Linux would be the Operating System name, although we all refer to it as "Linux"

```
flavio@flavio: ~
flavio@flavio: ~$ sw_vers
ProductName: Mac OS X
ProductVersion: 10.15.6
BuildVersion: 19G2021
flavio@flavio: ~
```

env

The `env` command can be used to pass environment variables without setting them on the outer environment (the current shell).

Suppose you want to run a Node.js app and set the `USER` variable to it.

You can run

```
env USER=flavio node app.js
```

and the `USER` environment variable will be accessible from the Node.js app via the `Node process.env` interface.

You can also run the command clearing all the environment variables already set, using the `-i` option:

```
env -i node app.js
```

In this case you will get an error saying `env: node: No such file or directory` because the `node` command is not reachable, as the `PATH` variable used by the shell to look up commands in the common paths is unset.

So you need to pass the full path to the `node` program:

```
env -i /usr/local/bin/node app.js
```

Try with a simple `app.js` file with this content:

```
console.log(process.env.NAME)
console.log(process.env.PATH)
```

You will see the output being

```
undefined
undefined
```

You can pass an env variable:

```
env -i NAME=flavio node app.js
```

and the output will be

```
flavio
undefined
```

Removing the `-i` option will make `PATH` available again inside the program:

```
flavio@flavio: ~
flavio@flavio: ~$ env NAME=flavio node app.js
flavio
/usr/local/bin:/usr/bin:/bin:/usr/sbin:/sbin:/Library/Apple/usr/bin
```

The `env` command can also be used to print out all the environment variables, if ran with no options:

```
env
```

it will return a list of the environment variables set, for example:

```
HOME=/Users/flavio
LOGNAME=flavio
PATH=/usr/local/bin:/usr/bin:/bin:/usr/sbin:/sbin:/l
PWD=/Users/flavio
SHELL=/usr/local/bin/fish
```

You can also make a variable inaccessible inside the program you run, using the `-u` option, for example this code removes the `HOME` variable from the command environment:

```
env -u HOME node app.js
```

printenv

A quick guide to the `printenv` command, used to print the values of environment variables

In any shell there are a good number of environment variables, set either by the system, or by your own shell scripts and configuration.

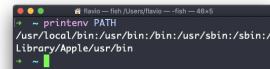
You can print them all to the terminal using the `printenv` command. The output will be something like this:

```
HOME=/Users/flavio
LOGNAME=flavio
PATH=/usr/local/bin:/usr/bin:/bin:/usr/sbin:/sbin:/l
PWD=/Users/flavio
SHELL=/usr/local/bin/fish
```

with a few more lines, usually.

You can append a variable name as a parameter, to only show that variable value:

```
printenv PATH
```



```
printenv PATH
/usr/local/bin:/usr/bin:/bin:/usr/sbin:/sbin:/Library/Apple/usr/bin
```

Conclusion

Thanks a lot for reading this book.

I hope it will inspire you to know more about Linux and its capabilities.

To learn more, check out my blog flaviocopes.com.

Send any feedback, errata or opinions at hey@flaviocopes.com